

Support Vector Machine

Nicolas Keriven
CNRS, IRISA, Rennes

(material from Florent Chatelain, Olivier Michel)

ENSTA 2023

Table of Contents

Support Vector Machine (SVM)

- Separating Hyperplane

- Separable case

- Nonseparable case

- Linear discrimination : comparison of SVM vs LDA

- Transformed space and Kernel function

- Examples

- Multiclass SVM

- SVM vs Logistic regression (LR)

Conclusions

Appendix : Some words on Random Forests

Support Vector Machine (SVM)

Theory elaborated in the early 1990's (Vapnik *et al*) based on the idea of 'maximum margin'

- ▶ deterministic criterion learned on the training set ← supervised classification
- 👉 general, i.e. model free, linear classification rule
- 👉 classification rule is linear in a transformed space of higher (possibly infinite) dimension than the original input feature/predictor space

Supplementary materials

- 📺 Coursera online video with python notebook material (13mn)
<https://www.coursera.org/lecture/data-analytics-accountancy-2/introduction-to-support-vector-machine-dDP0v>
- 🌐 Wikipedia page (quite complete and detailed)
https://en.wikipedia.org/wiki/Support_vector_machine
- 📄 Short and easy to understand Scikit-learn documentation (with examples)
<https://scikit-learn.org/stable/modules/svm.html>

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM

SVM vs Logistic regression (LR)

Conclusions

Appendix : Some words on Random Forests


Linear discrimination and Separating hyperplane

Binary classification problem

- ▶ $x \in \mathbb{R}^d$
- ▶ $y \in \{-1, 1\} \leftarrow 2 \text{ classes}$
- ▶ Training set (x_i, y_i) , for $i = 1, \dots, n$

Defining a **linear** discriminant function $h(x) \Leftrightarrow$ defining a separating **hyperplane** \mathcal{H} with equation

$$x^T \beta + \beta_0 = 0,$$

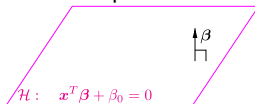


- ▶ $\beta \in \mathbb{R}^d$ is the normal vector (vector normal to the hyperplane \mathcal{H}),
- ▶ $\beta_0 \in \mathbb{R}$ is the intercept (regression interpretation) or offset (geometrical interpretation)
- 👁 \mathcal{H} is an *affine subspace* of dimension $p - 1$
- 👁 $h(x) \equiv x^T \beta + \beta_0$ is the associated (linear) discriminant function

Separating hyperplane and prediction rule

For a given separating hyperplane \mathcal{H} with equation

$$x^T \beta + \beta_0 = 0,$$



the prediction rule can

be expressed as

$$\hat{y} = \begin{cases} +1 & \text{if } h(x) = x^T \beta + \beta_0 \geq 0, (x \text{ is above } \mathcal{H}) \\ -1 & \text{otherwise, } (x \text{ is below } \mathcal{H}) \end{cases}$$

or in an equivalent way :

$$\hat{y} \equiv G(x) = \text{sign} [x^T \beta + \beta_0]$$

Rk : x is in class $y \in \{-1, 1\}$: prediction $G(x)$ is correct iff $y (x^T \beta + \beta_0) \geq 0$

Table of Contents

Support Vector Machine (SVM)

- Separating Hyperplane

- Separable case**

- Nonseparable case

- Linear discrimination : comparison of SVM vs LDA

- Transformed space and Kernel function

- Examples

- Multiclass SVM

- SVM vs Logistic regression (LR)

Conclusions


Appendix : Some words on Random Forests

Separating Hyperplane : separable case

Linear separability assumption : $\exists \beta \in \mathbb{R}^d$ and $\beta_0 \in \mathbb{R}$ s.t. the hyperplane $x^T \beta + \beta_0 = 0$ perfectly separates the two classes on the training set :

$$y_i (x_i^T \beta + \beta_0) \geq 0, \quad \text{for } i = 1, \dots, n,$$

Separable case ($p = 2$ example)



Pb : infinitely **many** possible perfect **separating hyperplanes** $x^T \beta + \beta_0 = 0$
Find the 'optimal' separating hyperplane ?
With the best

Maximum margin separating hyperplane (separable case)

Distance of a point x_k to an hyperplane \mathcal{H} s.t. $x^T \beta + \beta_0 = 0$,

$$d(x_k, \mathcal{H}) \equiv \min_x \left\{ \|x - x_k\| : x^T \beta + \beta_0 = 0 \right\}$$

Maximum margin principle

We are interested in the 'optimal' perfect separating hyperplane maximizing the distance $M > 0$, called the **margin**, between the samples of each class and the separating hyperplane

⇒ Find $\beta \in \mathbb{R}^d$ and $\beta_0 \in \mathbb{R}$ s.t. the margin

$$M = \min_{1 \leq k \leq n} \{d(x_k, \mathcal{H})\}$$


is **maximized**

Signed distance

From the orthogonality principle,

$$d(x_0, \mathcal{H}) = \|x_0 - \hat{x}_0\|,$$

where \hat{x}_0 is the orthogonal projection of x_0 on \mathcal{H} .



$\Rightarrow x_0 - \hat{x}_0$ and β are collinear,

$\Rightarrow x_0 - \hat{x}_0 = \underbrace{\langle x_0 - \hat{x}_0, \beta^* \rangle}_{\text{signed distance}} \beta^*$, where $\beta^* = \frac{\beta}{\|\beta\|}$,

$$\Rightarrow \text{signed distance} = (x_0 - \hat{x}_0)^T \frac{\beta}{\|\beta\|} = \frac{x_0^T \beta - \hat{x}_0^T \beta}{\|\beta\|} = \frac{x_0^T \beta + \beta_0}{\|\beta\|},$$

Remarks

- ▶ $|\langle x_0 - \hat{x}_0, \beta^* \rangle| = \|x_0 - \hat{x}_0\| = d(x_0, \mathcal{H}) \leftarrow$ “unsigned distance”
- ▶ for any perfect separating hyperplane $y_k \langle x_k - \hat{x}_k, \beta^* \rangle = \frac{1}{\|\beta\|} y_k (x_k^T \beta + \beta_0) \geq 0$, for $k = 1, \dots, n$: the signed distance must have the sign of y

Canonical separating hyperplane

For any perfect separating hyperplane, for $i = 1, \dots, n$


$$y_i \langle x_i - \hat{x}_i, \beta^* \rangle = d(x_i, \mathcal{H})$$

Hence, the margin reads

$$M \equiv \min_{1 \leq i \leq n} \{d(x_i, \mathcal{H})\} = \frac{1}{\|\beta\|} \min_{1 \leq i \leq n} \{y_i(x_i^T \beta + \beta_0)\}$$

- ▶ The bound M is reached (min of a finite set),
- ▶ the samples at the margin are denoted as x_{margin}
- ▶ Canonical expression of the separating hyperplane : remark that \mathcal{H} can be defined with β, β_0 up to a multiplicative constant. Hence β and β_0 are normalized s.t.

$$y_{\text{margin}}(x_{\text{margin}}^T \beta + \beta_0) = 1, \quad \text{thus } M = \frac{1}{\|\beta\|}$$



Primal problem (separable case)

Canonical hyperplane expression :

$$\begin{aligned} \text{maximizing the margin } M = \frac{1}{\|\beta\|} &\Leftrightarrow \text{minimizing } \|\beta\| \\ &\Leftrightarrow \text{minimizing } \frac{1}{2} \|\beta\|^2 \end{aligned}$$

Primal optimization problem

$$\begin{cases} \min_{\beta, \beta_0} & \frac{1}{2} \|\beta\|^2, \\ \text{subject to} & y_k (x_k^T \beta + \beta_0) \geq 1, \text{ for } 1 \leq k \leq n. \end{cases}$$

- ▶ quadratic criterion + linear inequality constraints
- 👉 convex optimization problem for which standard numerical procedures are available

Reminder on constrained optimization

Constrained problem : primal problem


$$\begin{cases} \min_x & f(x) \\ \text{s.t.} & g(x) \leq 0 \end{cases}$$

Objective function $f(x)$

To decrease the objective function $f(x)$, a descent direction \mathbf{d} must satisfy

$$f(x + \epsilon \mathbf{d}) \approx f(x) + \epsilon \nabla f(x)^T \mathbf{d} < f(x),$$

hence \mathbf{d} is a **descent direction** iff $\nabla f(x)^T \mathbf{d} < 0$



Reminder on constrained optimization

Objective $f(x)$

$$\text{descent direction : } \nabla f(x)^T \mathbf{d} < 0$$

Constraint $g(x)$

To satisfy the constraint, a **feasible** descent direction \mathbf{d} must satisfy

$$g(x + \epsilon \mathbf{d}) \approx g(x) + \epsilon \nabla g(x)^T \mathbf{d} \leq 0,$$


hence

$$\text{feasible direction : } \begin{cases} g(x) < 0 & \Rightarrow \text{no constraint on } \mathbf{d}, \\ g(x) = 0 & \Rightarrow \nabla g(x)^T \mathbf{d} \leq 0 \end{cases}$$

Reminder on constrained optimization (Cont'd)

Necessary conditions : two possibilities for optimality. There is no feasible descent direction in x^* when either

1. $g(x^*) < 0$ and the minimum is reached
 $\nabla f(x^*) = 0$: same condition as unconstrained
2. $g(x^*) = 0$ and $\nabla f(x^*), \nabla g(x^*)$ are in **opposite direction** : $\nabla f(x^*) = -\alpha \nabla g(x^*)$ with $\alpha > 0$ (the constraints are **saturated**)



Reminder on constrained optimization (Cont'd)

Constrained form : primal problem

$$\begin{cases} \min_x & f(x) \\ \text{s.t.} & g_j(x) \leq 0, \text{ for all } j = 1, \dots, q \end{cases}$$

Lagrangian form : dual problem

Inequality convex constraints \Rightarrow introduction of the **Lagrange multipliers** α_j

$$\mathcal{L}(x, \alpha) = f(x) + \sum_j \alpha_j g_j(x)$$

A **saddle point** (x^*, α^*) , that is such that $\mathcal{L}(x^*, \alpha) \leq \mathcal{L}(x^*, \alpha^*) \leq \mathcal{L}(x, \alpha)$, with $\alpha_j^* \geq 0$, yields an optimal x^* . In this case, $\min_x \max_\alpha \mathcal{L} = \max_\alpha \min_x \mathcal{L}$.

Karush-Kuhn-Tucker (KKT) conditions

For x^* being a local min, it is **necessary** that (generally sufficient...)

$$\begin{cases} \nabla f(x^*) + \sum_{j=1}^q \alpha_j \nabla g_j(x^*) = 0 & \leftarrow \text{first order conditions} \\ \text{s.t. } \alpha_j \geq 0 \text{ and } \alpha_j g_j(x^*) = 0 & \leftarrow \text{complementary conditions} \end{cases}$$

Lagrangian for SVM (separable case)

Linear constraints of positivity \Rightarrow introduction of the Lagrange multipliers

Lagrangian

$$L(\beta, \beta_0, \alpha) = \frac{1}{2} \|\beta\|^2 - \sum_{i=1}^n \alpha_i \underbrace{\left[y_i (x_i^T \beta + \beta_0) - 1 \right]}_{\geq 0},$$

where α_i are the Lagrange multipliers

First order Karush–Kuhn–Tucker necessary conditions

Setting the partial derivatives w.r.t. β and β_0 to zero yields

$$\begin{cases} \hat{\beta} &= \sum_{i=1}^n \alpha_i y_i x_i, \\ 0 &= \sum_{i=1}^n \alpha_i y_i, \end{cases}$$

► plugging these expression in the Lagrangian yields the dual expression

Dual problem (separable case)

Other mean of getting the solution (with different properties...)

Dual optimization problem

$$\begin{cases} \max_{\alpha} & \tilde{L}(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j, \\ \text{subject to} & \alpha_i \geq 0 \text{ and } \sum_{i=1}^n \alpha_i y_i = 0. \end{cases}$$

- simple convex optimization problem for which standard numerical procedures are available
- calculation of the optimum multipliers $\hat{\alpha}_i$
- then $\hat{\beta} = \sum_i \hat{\alpha}_i y_i x_i$ (see after for β_0)

Support vectors and maximum margin hyperplane (separable case)

Complementary slackness Karush-Kuhn-Tucker necessary conditions

With $h(x) = \beta^T x + \beta_0$:

$$\hat{\alpha}_i [y_i h(x_i) - 1] = 0 \quad \Rightarrow \quad \hat{\alpha}_i = 0 \quad \text{as} \quad y_i h(x_i) > 1$$

- ▶ either $\hat{\alpha}_i = 0$, or $h(x_i) = y_i$ and x_i is at the margin
- ▶ since $\hat{\beta} = \sum_{i=1}^n \hat{\alpha}_i y_i x_i$, $\hat{\beta}$ depends only on the points at the margin, aka **support vectors**
- ▶ $\hat{\beta}_0$ can be derived from *any* of support vectors x_{margin} (ie for which $\hat{\alpha}_i \neq 0$) :

$$\hat{\beta}^T x_{\text{margin}} + \hat{\beta}_0 = y_{\text{margin}} \Rightarrow \hat{\beta}_0 = -\hat{\beta}^T x_{\text{margin}} + y_{\text{margin}}$$


- ☞ the only **inputs used to construct the maximum margin hyperplane** are the **support vectors** and the discriminant function reads

$$h(x) = \sum_{i=1}^n \hat{\alpha}_i y_i (x - x_{\text{margin}})^T x_i + y_{\text{margin}}$$

Maximum margin separating hyperplane (separable case)

Separable case

- ➦ Maximizing the *margin* M between the separating hyperplane and the training data :



The maximum margin hyperplane depends only on the points at the margin called the *support vectors*

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM


SVM vs Logistic regression (LR)

Conclusions

Appendix : Some words on Random Forests

Nonseparable case

- ▶ in general, overlap of the 2 classes
- 👉 No hyperplane that perfectly separates the training data



Maximum margin separating hyperplane (nonseparable case)

Soft-Margin solution for the nonseparable case

Considering a **soft-margin** that allows wrong classifications

- ▶ introduction of *slack variables* $\xi_i \geq 0$ s.t.

$$y_i(x_i^T \beta + \beta_0) \geq 1 - \xi_i$$

Support vectors include now the wrong classified points, and the points inside the margins ($\xi_i > 0$)

- ▶ Primal problem : adding a penalty in the criterion

$$\begin{cases} \min_{\beta, \beta_0, \xi} & \frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^n \xi_i, \\ \text{subject to} & y_i(x_i^T \beta + \beta_0) \geq 1 - \xi_i, \xi_i \geq 0 \end{cases}$$

where $C > 0$ is the “cost” or “regularization” parameter.

- ▶ Class `LinearSVC` in Scikit-learn, main hyper-parameter is C

Rk : the optimal value of ξ_i is $L(h(x_i), y_i) = \max(0, 1 - y_i h(x_i))$, aka the **hinge loss**.

Regularization parameter (nonseparable case)

$$\text{Criterion to be minimized : } \frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^n \xi_i,$$

Influence of the regularization parameter $C > 0$

C drives the margin size, thus the number of support vectors

- ▶ $C \gg 0$: **small margin**, less support vectors (\sim overfitting) $C \rightarrow +\infty$: converges in the separable case to the *Hard-Margin* solution
- ▶ $C \rightarrow 0^+$: **large margin**, more support vectors (\sim underfitting)

Rk : strength of the regularization is inversely proportional to C (compared with the regularization parameter λ for ridge penalty, $C \equiv \frac{1}{\lambda}$)

Choosing the regularization parameter $C > 0$

- ▶ as usual, the optimal C can be estimated by cross validation
- 👉 performance might not be very sensitive to choices of C (due to the rigidity of a linear boundary)
- 👉 usually $C \approx 1$ yields a good trade-off

Dual problem (nonseparable case)

Introducing the Lagrangian and substituting the first order KKT conditions w.r.t. β , β_0 , ξ yields the dual expression

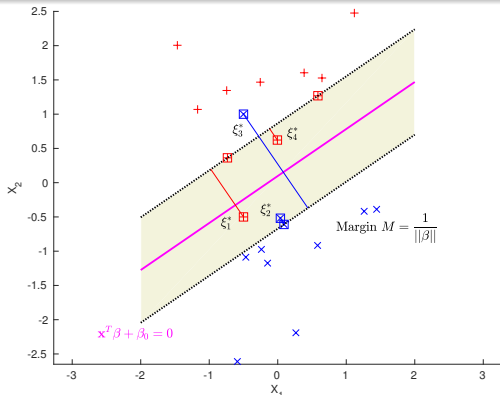
Dual optimization problem

$$\begin{cases} \max_{\alpha} & \tilde{L}(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j, \\ \text{subject to} & 0 \leq \alpha_i \leq C \text{ and } \sum_{i=1}^n \alpha_i y_i = 0. \end{cases}$$

- only difference w.r.t the separable case : $\alpha_i \leq C$ constraint !
- sample on the margin are those for which $0 < \hat{\alpha}_i < C$ (strictly !)
- $\hat{\beta}, \hat{\beta}_0$ can be recovered from $\hat{\alpha}_i$ in the same manner as the hard-margin case
- simple convex optimization problem for which standard numerical procedure are available

Optimal separating hyperplane

Soft-Margin example (nonseparable case)



Vector Supports

The support vectors are now the points at the margin, inside the margin, or wrongly classified.

$\xi_i^* \equiv M \xi_i \leftarrow$ distance between a support vector and the margin

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM

SVM vs Logistic regression (LR)


Conclusions

Appendix : Some words on Random Forests


Linear discrimination : SVM vs LDA

Linear discrimination

- ▶ Linear Discriminant Analysis (LDA) : Gaussian generative model
- ▶ SVM : criterion optimization (maximizing the margin)




SVM




LDA

Linear discrimination : SVM vs LDA (Cont'd)

Adding one atypical data far from the others :



SVM



LDA

SVM property

- ▶ SVM is **insensitive** to outliers that are **far from the margin** (since they are not support vectors)
- ▶ SVM is sensitive to data that is **close to the margin** (esp. when C is large), they may change the set of support vectors.

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM

SVM vs Logistic regression (LR)

Conclusions


Appendix : Some words on Random Forests

A scatter plot showing two classes of data points in a 2D space. The horizontal axis is labeled X_1 and ranges from -2.5 to 2.0. The vertical axis is labeled X_2 and ranges from -2.0 to 1.5. The first class, represented by blue 'x' marks, consists of 15 points. The second class, represented by red '+' marks, consists of 15 points. The blue points are generally located on the periphery of the plot, while the red points are clustered in the central region.

31 / 54

Nonlinear discrimination in the input space

► $X \in \mathbb{R}^2$, $\phi(x) = (x_1^2, x_2^2, \sqrt{2}x_1x_2)^T$



Linear separation in the feature space $\mathcal{F} \Rightarrow$ Nonlinear separation in the input space

Kernel trick

Recall the dual formulation of SVM :

$$\begin{cases} \max_{\alpha} & \tilde{L}(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j, \\ \text{subject to} & 0 \leq \alpha_i \leq C \text{ and } \sum_{i=1}^n \alpha_i y_i = 0. \end{cases}$$

and the prediction function, for any $x_{\text{margin}} = x_i$ for which $0 < \hat{\alpha}_i < C$:

$$h(x) = \sum_{i=1}^n \hat{\alpha}_i y_i (x - x_{\text{margin}})^T x_i + y_{\text{margin}}$$

- ▶ Both only depends on the inner products $\langle x_i, x_j \rangle$ and $\langle x_i, x \rangle$
- ▶ Thus we can apply the kernel trick as previously for PCA or ridge regression.
- ▶ Class SVC in Scikit-learn uses a Gaussian kernel by default !

Reminder on Kernel trick

Kernel trick

Use of a kernel function k associated with an expansion/feature map ϕ :

$$\begin{aligned} k : \mathbb{R}^d \times \mathbb{R}^d &\rightarrow \mathbb{R} \\ (x, x') &\mapsto k(x, x') \equiv \langle \phi(x), \phi(x') \rangle_{\mathcal{H}} \end{aligned}$$

- ▶ explicit representations of the feature map ϕ and enlarged feature space \mathcal{H} are not necessary, only the expression of k is required
- ▶ regular SVM for linear kernel $k(x, x') = x^\top x'$
- ▶ The cost of dual SVM is $O(C_k n^2)$ (where C_k is the cost of computing k , generally $O(d)$), while the **primal** linear SVM was $O(nd^2 + d^3)$. Primal is generally faster, but the kernel trick is only possible in the dual.

Some properties of Kernel function

Definition (Positive semi-definite kernel)

A **symmetric** kernel $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$ is said to be positive semi-definite (psd) iff

$$\forall n \in \mathbb{N}, \quad \forall \xi_1 \dots \xi_n \in \mathbb{R}, \quad \forall \mathbf{x}_1 \dots \mathbf{x}_n \in \mathbb{R}^d, \quad \sum_{i,j} \xi_i \xi_j k(\mathbf{x}_i, \mathbf{x}_j) \geq 0$$

This is true for any inner product : $\sum_{i,j} \xi_i \xi_j \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle_{\mathcal{H}} = \| \sum_i \xi_i \phi(\mathbf{x}_i) \|_{\mathcal{H}}^2$.

But the converse is also true !!

Theorem (Mercer Theorem)

For every positive semi-definite kernel k , there exists a Hilbert space \mathcal{H} and a feature map $\phi : \mathbb{R}^d \rightarrow \mathcal{H}$ such that $k(\mathbf{x}, \mathbf{x}') = \langle \phi(\mathbf{x}), \phi(\mathbf{x}') \rangle_{\mathcal{H}}$.

It is possible to choose \mathcal{H} as a **space of functions**, such that $k(\cdot, \mathbf{x}) \in \mathcal{H}$, and $\langle f, k(\cdot, \mathbf{x}) \rangle_{\mathcal{H}} = f(\mathbf{x})$. It is called the **Reproducing Kernel Hilbert Space** (RKHS) associated to k .

Operations on kernels

Let k_1 and k_2 be psd, and $\lambda_{1,2} > 0$ then :

1. $\lambda_1 k_1$, (multiplication by a positive scalar)
2. $\lambda_1 k_1 + \lambda_2 k_2$, (sum of kernels),
3. $k_1 k_2$, (product of kernels),
4. $\exp(k_1)$, (exponential of kernel),
5. $(x_i, x_j) \mapsto g(x_i)g(x_j)k_1(x_i, x_j)$, with $g : \mathbb{R}^d \rightarrow \mathbb{R}$, (multiplication by a function)

are all positive semi-definite, hence **valid kernels**.

- 🔑 These operations allow us to create more complicated kernels by combining simple ones.

Choosing the Kernel function

Usual kernel functions

- ▶ Linear kernel ($\mathcal{F} \equiv \mathbb{R}^d$) : $k(x, x') = x^T x'$
- ▶ Polynomial kernel (dimension of \mathcal{F} increases with the order d)

$$k(x, x') = (x^T x')^d \quad \text{or} \quad (x^T x' + 1)^d$$

- ▶ Gaussian radial function (\mathcal{F} with infinite dimension)

$$k(x, x') = \exp(-\gamma \|x - x'\|^2)$$

- ▶ Neural net kernel (\mathcal{F} with infinite dimension)

$$k(x, x') = \tanh(\kappa_1 x^T x' + \kappa_2)$$


- ▶  standard practice is to estimate the optimal kernel parameters by **cross-validation**

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM


SVM vs Logistic regression (LR)

Conclusions

Appendix : Some words on Random Forests

Application : binary data

Linear kernel




$$C = 10000$$

- SVM decision boundary
- - - SVM margin boundaries
- - - Bayes (optimal) decision boundary

Application : binary data

Polynomial kernel ($d = 4$)




$C \approx 1$

- SVM decision boundary
- SVM margin boundaries
- Bayes (optimal) decision boundary

Application : binary data

Gaussian radial kernel ($\gamma = 1$)



- SVM decision boundary
- SVM margin boundaries
- - - Bayes (optimal) decision boundary

$C \approx 1$

Scale your data !

Scaling of the variables matters !

For instance, with Gaussian kernel

$$k(x, x') = \exp(-\gamma \|x - x'\|^2) = \exp\left(-\gamma \sum_{i=1}^p (x_i - x'_i)^2\right),$$

the variables that have the greatest magnitudes are favored to compute distances or inner-products.

Practical advices

- ▶ If the variables are in different units, scaling each is **strongly recommended**.
- ▶ If they are in the same units, you might or might not scale the variables (depend on your problem)

Usual scaling methods

- ▶ **normalization** in $[0, 1]$: $\tilde{x}_i = \frac{x_i - \min_i}{\max_i - \min_i}$
- ▶ **standardization** to get zero mean and unit variance : $\tilde{x}_i = \frac{x_i - \mu_i}{\sigma_i}$

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM

SVM vs Logistic regression (LR)

Conclusions

Appendix : Some words on Random Forests

Multiclass SVM

- ▶ $Y \in \{1, \dots, K\} \leftarrow K$ classes

Standard approach : direct generalization by using **multiple binary SVMs**

OVA : one-versus-all strategy

- ▶ K classifiers between one class (+1 label) versus all the other classes (-1 label)
- 👉 classifier with the highest confidence value (e.g. the maximum distance to the separator hyperplane) assigns the class

OVO : one-versus-one strategy

- ▶ $\binom{K}{2} = K(K-1)/2$ classifiers between every pair of classes
- 👉 majority vote rule : the class with the most votes determines the instance classification

Which to choose? if K is not too large, choose OVO

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM

SVM vs Logistic regression (LR)

Conclusions

Appendix : Some words on Random Forests

SVM vs Logistic regression (LR)

- ▶ When classes are nearly separable, SVM does better than LR. So does LDA.
- ▶ When not, LR (with ridge penalty) and SVM are very similar
- ▶ If one wants to **estimate probabilities** for each class, LR is the natural choice
- ▶ For **non linear boudaries**, kernel SVMs are popular. Can use kernels with LR and LDA as well, but computations are more expensive.

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM

SVM vs Logistic regression (LR)

Conclusions

Appendix : Some words on Random Forests

Conclusions on Support Vector Machines

- ▶ **model free approach** based on a maximum margin criterion : may be very efficient for real-world data (but do not directly provide probability estimates nor variable importance weights)
- ▶ **memory efficient** sparse solution characterized by the only support vectors
- ▶ **versatile** algorithm : different choices of kernels to make a **nonlinear classification** in the original input space by performing an implicit linear classification in a higher dimensional space
- ▶ Possible **extensions** to other tasks than classification like **regression** (*support vector regression*) or **anomaly detection** (*one-class SVM*)
- ▶ **effective in high dimensional** spaces even when $p > n$.
- ▶ **computationally expensive** to train for large n data sets : cost of the optimization procedure to solve the quadratic problem scales from $O(pn^2)$ to $O(pn^3)$ operations depending on the training set.
- ▶ **popular algorithm**, with a large literature

Perspectives on 'Black Box' (model free) approaches

Random Forests (not in this course)

- ▶ involve **decision trees** to split the prediction space in simple regions
- ▶ **combine** multiple decision trees to yield a single consensus prediction
- 👉 method able to scale efficiently to high dimensional data and large data sets

Deep Neural Nets

- ▶ Neural Nets with multiple hidden layers between input and output ones
- ▶ many variants of deep architectures (Recurrent, Convolutional,...) used in specific domains (speech, vision, ...)
- ▶ very computationally expensive to train due to the high number of parameters
- ▶ supported by empirical evidence
- 👉 dramatic performance jump for some big data applications

Table of Contents

Support Vector Machine (SVM)

Separating Hyperplane

Separable case

Nonseparable case

Linear discrimination : comparison of SVM vs LDA

Transformed space and Kernel function

Examples

Multiclass SVM

SVM vs Logistic regression (LR)





Conclusions

Appendix : Some words on Random Forests

Random Forests




- ▶ Introduced in 2001 (Breiman)
- ▶ Model free and non linear
- ▶ Build a large collection of de-correlated trees and average them
- ▶ Combination of weak learners

Decision trees



Random Forests

- ▶ For each tree :
 - ▶ Draw bootstrap sample X^b for training sample
 - ▶ Learn tree, for each node
 - ▶ select m features from the initial p features
 - ▶ Find the best split (e.g. Gini index, entropy ...)



Application : binary data

